

SIDDHARTH GOEL

New York City, NY | 413-695-8052 | siddharth.goel7@gmail.com | [linkedin.com/in/sidgoel7](https://www.linkedin.com/in/sidgoel7)

EDUCATION

University of Massachusetts Amherst, Amherst, MA

09-01-2021 - 05-22-2023

Master of Science in Computer Science, GPA: 3.84/4.0

Coursework: Neural Networks, Systems for Data Science, Secure Distributed Systems, Machine Learning, Distributed and Operating Systems, Game Programming, Algorithms for Data Science, Intro to Data Science

Thapar Institute of Engineering and Technology, Punjab, IN

07-26-2015 - 05-30-2019

Bachelor of Engineering in Computer Engineering, CGPA: 9.32/10

Coursework: Data Structures and Algorithms, Operating Systems, Image Processing, Secure Coding, Computer Networks

SKILLS

Programming/Scripting Languages: Java [Advanced], Scala [Advanced], Python [Intermediate], Rust [Intermediate], SQL

Tools/Platform: Linux - CentOS, Unix, Git / Mercurial, Vim, IntelliJ, VSCode, Jira, Agile, Maven, Buck, Airflow

Frameworks and Technologies: Spark, Hadoop, Kafka, NumPy, Redis, Hive, Iceberg, Ansible, gRPC

EXPERIENCE

Software Engineer III - Walmart Global Tech, Hoboken, NJ

06-17-2024 - Current

- Working in the AdTech Measurement & Reporting Team to develop a generic ad attribution platform (**Attribution As A Service**) which attributes transactions to advertisement activities based on various attribution algorithms using Spark, Scala, Airflow and Python
- The attribution platform will make use of its generic capabilities to reduce development time of new attribution pipelines
- Developed the Path-to-Conversion reporting module for the Walmart Connect business. This Module shows reports for Path-to-Conversion attributing sales to the ads published by the advertisers.
- Contributing to the department's Operational Excellence endeavors by being the OE Champion in my team
- Technologies used: Apache Spark, Apache Airflow, Google Cloud Platform, Scala, Python
- Skills used: Cloud computing (GCP - DataProc), Distributed Systems, Software Engineering, Data Analytics

Senior Software Engineer - Goldman Sachs, New York City, NY

07-23-2023 - 06-14-2024

- Working in the Compliance Big Data Engineering team to develop high performance data pipelines which detect exceptions in market trading data
- Was responsible for **performance improvement** of the Goldman Sachs Cross Trade tool. Reduced trade submission times from 2 minutes to 10 seconds per transaction by using **distributed processing** techniques
- Responsible for design, development and successful deployment of the team's new **core data pipeline** in Spark

Graduate Student Researcher - IBM, Remote

01-24-2023 - 05-22-2023

- Experimented on the best ways to exploit representations of polymers stemming from different pre-trained models
- Implemented the research findings on polymer property prediction tasks, improving accuracy and performance by 5%

Software Engineer Intern - Meta Platforms Inc. (formerly Facebook), Menlo Park, CA

05-22-2022 - 08-12-2022

- Developed **Hargow** – a fast command line tool which facilitates in execution of complex workloads in production environments of choice – using Rust
- Implemented features such as **metadata store** and **remote job execution** as part of the project
- Integrated the **Citrix** remote user desktop with the **hargow-server** using **rust curl-lib** and **thrift**

Software Engineer Intern - Bytedance Inc., New York City, NY

01-31-2022 - 04-22-2022

- Improve ingestion performance of Bytedance's Offline Inference Framework by 10% by optimizing the (de)compression logic in their pipeline as part of the Bytedance Infrastructure Lab.
- Contributed to the open source [Ray project](#) – Added support for decompression of snappy compressed txt files in Ray.

Software Development Engineer (Big Data) - Guavus Network Systems, Gurgaon, IN

07-15-2019 - 07-16-2021

- Designed and developed a Tethering detection solution which analyzes EDR datasets by performing **deep packet inspection** and applying filtering rules to identify tethered smart devices in a cellular network
- Created multiple micro-applications which analyze LTE data from CSPs to derive **customer experience** and identify potential pain points in the network using **Anomaly Detection**

- Developed a **streaming service** which processes data from **Kafka** to maintain a live **Redis** database of IP and mobile number mappings, with **latency < 0.5s**
- Handled team's maintenance activities in rotation, fixing **4 critical** and **20+ non-critical** production issues